

# Report

## **Assignment 1 Open Source Firewalls Jan 2024**

Student Name  
**Sebastian Konefal**

Student Number  
**b00168561**

email address  
**b00168561@mytudublin.ie**

TU765  
Digital Forensics & Cyber Security

DFCS H3015 - Network Security [1114.VLE-BN]

Date

03/03/2024

## Table of Contents

Assignment 1 Open Source Firewalls Jan 2024.....	1
Introduction.....	3
1. NFtables development.....	4
1.1. Topology .....	4
1.2. My NFTable.....	6
1.3. Adjusting software and environment .....	7
1.3.1. Global settings on all VMs: .....	7
2.3.2. VM1-Server settings .....	9
2.3.3. VM2- Firewall settings .....	10
2.3.4. VM3 - Attacker settings.....	10
3. Testing.....	11
Summary of tests .....	11
3.1. Open port scan on VM1-Server .....	12
3.2. Open port scan on VM1-Server – NFTable off .....	14
3.3. Open port scan on VM-Firewall .....	14
3.4. Open port scan on VM2-Firewall – NFTable off.....	15
3.5. IPv6 open port scan on VM-Firewall.....	15
3.6. IPv6 open port scan on VM2-Server.....	17
3.7. Accessing Web Server on VM1-Server .....	17
3.8. DOS on VM2-Server port 80 .....	18
3.9. DOS on VM2-Server port 80 – NFTable off .....	19
3.10. DOS on VM2-Firewall intnet on port 22 using IP of non-privileged range .....	19
3.11. DOS on VM2-Firewall on port 22 using IP of privileged range - IP 192.168.5.30 used .....	21
3.12. FIXED - DOS on VM2-Firewall on port 22 using IP of privileged range - IP 192.168.5.30 used .....	22
3.13. DOS on VM2-Firewall on port 22 using IP of privileged range - IP 192.168.5.30 used with Nftable off.....	24
3.14. Triggering egress hook by accessing web server .....	26
3.15. Connecting to FTP (hidden open port) & SSH .....	26
4. Conclusions.....	27
5. Source editable NFTable:.....	28
6. References:.....	33

## Introduction

Nftables is a subsystem of the Linux kernel that provides a set of functions for filtering and network manipulation packets. Nftables replaced the packet filtering system - iptables that was used by Linux for many years. Compared to iptables, nftables is engineered to use resources more efficiently, it can manage larger volumes of network traffic while putting a lower load on the system. Creating and maintaining packet filtering rules is made simpler by more flexible and easy-to-understand configuration syntax.<sup>1</sup> There are no pre-defined tables or chains when using nftables. Each table needs to be defined and only consists of the objects (chains, sets, maps, flowtables and stateful objects). The nftables rule is composed of zero or more expressions followed by one or more statements. Linear evaluation from left to right applies to multiple expressions - every expression is evaluated one after another. When we get to the final expression, the statements in the rule are executed as the packet matches all of the expressions in the rule. Linear evaluation from left to right also applies to multiple statements - this means that a single rule can take multiple actions by using multiple statements.<sup>2</sup>

Nftables has been available since Linux kernel 3.13. It supports a variety of address families such as ip - IPv4, ip6 - IPv6, arp - Address Resolution Protocol (ARP), bridge - processing for bridged packets and netdev - Netdev address family. The address family to which the rule will apply is determined by the nftables processing architecture overall. Nftables makes use of one or more tables, each of which has chains that carry the processing rules. Those processing rules consist of statements like drop, queue, and continue and they are created by expressions such as the address, interface, ports, or other data in the packet. In Linux, certain address families have hooks that allow nftables to inspect packets as they traverse the network stack.<sup>3</sup> Hooks are triggered when incoming or outgoing packet enters a node with active nftables. Rules linked to these hooks are allowed to interact with network traffic by the Linux kernel. Netfilter contains five hooks including prerouting, input, output, postrouting, forward, and ingress.<sup>4</sup>




Fig. 1 Easillyy (n.d.) A comprehensive guide to nftables [online graphic]

The main functions of nftables are:

- packet filtering - Nftables enables to set up rules to filter network packets according to parameters including IP addresses, ports, protocols, and network interfaces, allowing the management of packets that can enter, exit, or pass through the system.

---

<sup>1</sup> Shirvar A., (2020)

<sup>2</sup> Wiki, (2021)


<sup>3</sup> Suehring S., 2015, 83-93

<sup>4</sup> Easillyy, (n.d.)

- Network Address Translation (NAT) - Network address translation (NAT) is a function of nftables that allows change of the IP addresses and ports of network packets which can be used for port translation (masquerading) and static network address translation (SNAT and DNAT).
- load balancing - It can improve the availability and scalability of network services by dividing incoming traffic across several backend servers providing load balancing.
- packet classification (QoS) - By implementing Quality of Service (QoS), nftables ensures adequate performance by prioritisation of some traffic types above others.
- packet logging - Nftables allows tracking traffic traversing the system and logging important events therefore it is very useful for network troubleshooting, auditing and monitoring.
- packet manipulation - Advanced packet manipulation features enable the modification of particular fields in network packets (such as IP addresses, ports, VLAN tags, and others) as they traverse the system.<sup>5</sup>

## 1. Nftables development

### 1.1. Topology



Installing xUbuntu on each VM, example below:

Who are you?

Your name: VM1-ServerUser ✓

Your computer's name: VM1-Server ✓  
The name it uses when it talks to other computers.

Pick a username: vm1-serveruser ✓

Choose a password: ●●●●●●●● Good password





Confirm your password: ●●●●●●●● ✓

☒ Log in automatically  
☐ Require my password to log in

Back Continue

<sup>5</sup> SW Team, (2024)

Created 3 VMs in total:

	<b>VM3-Attacker</b> ➡ Running	
	<b>VM1-Server</b> ➡ Running	
	<b>VM2-Firewall</b> ➡ Running	

## 1.2. My NFTable

### Editable text at the end of the report

**Possible further development:** Honeypots, Logging and Analysis, Geo-IP Filtering, Load Balancing

```
flush ruleset
table ip Server{

#Variables:
#Declouling list of IPs allowed to connect vis SSH
set allowed_ssh_ips{
    typeof ip saddr . tcp dport
    flags interval,constant #to use CIDR range and prevent changes from cl
    auto-merge #merge any overlaping range
    elements = {192.168.5.30/30 . 22}
}

#Counter variables
counter counter_ct_web{
}
counter counter_ct_ssh{
}
#3x timeout variables
set timeout1{
    typeof ip saddr
    flags timeout
}
set timeout2{
    typeof ip saddr
    flags timeout
}
set timeout3{
    typeof ip saddr
    flags timeout
}

#Allowed server ports
set allowed_server_ports{
    typeof tcp dport
    elements = {22, 80, 443}
}

#Variable to prevent flood attack on server
set frequent_server_sources {
    typeof ip saddr
    flags timeout
}

#variable to prevent flood attack on firewall
set frequent_firewall_sources {
    typeof ip saddr
    flags timeout
}

#CHAINS:
chain prerouting_server{
    type nat hook prerouting priority 0;
    #Rule to forward only specified IP - implemented for testing purposes
    iifname "enp0s8" ip saddr 192.168.5.30 tcp dport {80, 443 } counter dnat to 10.5.5.30:80-443
    #continue to statefull packet examination
}

chain server_forward {
    type filter hook forward priority 0;
    # Permit established and related SSH connections
    ct state established, related accept

    #Limiting new flood attack with 1s timeout
    ct state new ip saddr @frequent_server_sources counter drop
    ct state new add @frequent_server_sources {ip saddr timeout 1s}
    #Forawrding web services
    ct state new tcp dport {80, 443 }counter name counter_ct_web accept
    #Accepting icmp
    ip protocol icmp accept
    #Dropping malicious/invalid packets
    ct state invalid counter drop

    #JUMP to chain dealing with SSH port 22
    tcp dport 22 jump input_ssh

    #Dropping any other traffic
    counter drop
}
```

```

chain input_firewall {
    type filter hook input priority filter; policy accept;
    iifname lo counter accept
    #if a contract receives a SYN packet it considers connection as new , SYN-ACK considers as established
    ct state established, related counter accept

    #Limiting flood attacks with 1s timeout
    ct state new ip saddr @frequent_firewall_sources counter drop
    ct state new add @frequent_firewall_sources {ip saddr timeout 1s}

    #Accepting icmp
    ip protocol icmp accept
    #Dropping malicious/invalid packets
    ct state invalid counter drop

    #JUMP to chain dealing with SSH port 22
    tcp dport 22 jump input_ssh

    #Dropping any other traffic
    counter drop
}
chain input_ssh{
    # Permit established and related SSH connections
    ct state established, related accept
    #Dedicated rule to management access via allowed range of IPs -it will accept unlimited connections but not more often
    #then every 1s as per limiting flood attack rule in input_firewall chain
    ct state new ip saddr . tcp dport @allowed_ssh_ips counter accept #to prevent blocking management access
    #Time-outs for SSH
    ct state new ip saddr @timeout2 tcp dport 22 add @timeout3 {ip saddr timeout 3d}
    ct state new ip saddr @timeout1 tcp dport 22 add @timeout2 {ip saddr timeout 3m}
    ct state new tcp dport 22 add @timeout1 {ip saddr timeout 1m}
    ct state new ip saddr @timeout3 tcp dport 22 counter drop
    #Only SYN packet will match this rule
    ct state new tcp dport 22 counter name counter_ct_ssh accept
}

#OUTPUT
#variable to track outgoing connections:
set egress tcp_connections{
    typeof meta skuid . ip daddr . tcp dport;
    counter;
}
chain egress {
    type filter hook output priority filter; policy accept;
    ct state new ip protocol tcp add @egress_tcp_connections {meta skuid . ip daddr . tcp dport}
}
}

```

6789

## 1.3. Adjusting software and environment

### 1.3.1. Global settings on all VMs:

Installed Guest Additions for better user experience:

```

.../media/vm2-firewalluser/VBox_GAs_7.0.12$ ./autorun.sh

```

<sup>6</sup> LinuxCloudHacks (2023) Mastering Nftables Sets: A Comprehensive Guide

<sup>7</sup> LinuxCloudHacks (2023) Protecting Incoming Traffic with Nftables

<sup>8</sup> LinuxCloudHacks (2023) Unleashing the Power of Nftables Chains and Verdict Maps

<sup>9</sup> Wiki (2018)

```

Verifying archive integrity... 100% MD5 checksums are OK. All good.
Uncompressing VirtualBox 7.0.12 Guest Additions for Linux 100%
VirtualBox Guest Additions installer
Copying additional installer modules ...
Installing additional modules ...
VirtualBox Guest Additions: Starting.
VirtualBox Guest Additions: Setting up modules
VirtualBox Guest Additions: Building the VirtualBox Guest Additions kernel
modules. This may take a while.
VirtualBox Guest Additions: To build modules for other installed kernels, run
VirtualBox Guest Additions: /sbin/rcvboxadd quicksetup <version>
VirtualBox Guest Additions: or
VirtualBox Guest Additions: /sbin/rcvboxadd quicksetup all
VirtualBox Guest Additions: Building the modules for kernel 6.5.0-21-generic.

VirtualBox Guest Additions: Look at /var/log/vboxadd-setup.log to find out what
went wrong
VirtualBox Guest Additions: Running kernel modules will not be replaced until
the system is restarted or 'rcvboxadd reload' triggered
VirtualBox Guest Additions: reloading kernel modules and services
VirtualBox Guest Additions: kernel modules were not reloaded
VirtualBox Guest Additions: kernel modules and services were not reloaded
The log file /var/log/vboxadd-setup.log may contain further information.
Press Return to close this window...

```

Updating and upgrading using additional NAT interface:

```
~# apt-get update
```

```
~# apt-get upgrade
```

Installing Wireshark:

```
root@VM1-Server:~# apt install software-properties-common
```

```
root@VM1-Server:~# add-apt-repository ppa:wireshark-dev/stable
```

```
root@VM1-Server:~# apt-get install wireshark 10
```

Adding user to sudoers:

```
root@VM1-Server:~# sudo visudo
```

```

# User privilege specification
root    ALL=(ALL:ALL) ALL
vm1-serveruser ALL=(ALL:ALL) ALL 11

```

```
root@VM2-Firewall:~# apt install net-tools
```

---

<sup>10</sup> ZacsTech (2022)

<sup>11</sup> ZacsTech (2022)



### 2.3.2. VM1-Server settings

```
root@VM1-Server:~# ip a
1: lo: <LOOPBACK,UP,LOWER_UP> mtu 65536 qdisc noqueue state UNKNOWN group default qlen 1000
    link/loopback 00:00:00:00:00:00 brd 00:00:00:00:00:00
    inet 127.0.0.1/8 scope host lo
        valid_lft forever preferred_lft forever
    inet6 ::1/128 scope host
        valid_lft forever preferred_lft forever
2: enp0s3: <BROADCAST,MULTICAST,UP,LOWER_UP> mtu 1500 qdisc fq_codel state UP group default qlen 1000
    link/ether 08:00:27:0e:19:d8 brd ff:ff:ff:ff:ff:ff
    inet 10.5.5.30/24 brd 10.5.5.255 scope global noprefixroute enp0s3
        valid_lft forever preferred_lft forever
    inet6 fe80::43eb:1933:6699:4ab2/64 scope link noprefixroute
        valid_lft forever preferred_lft forever
```

Installing services:

```
root@VM1-Server:~# sudo apt install apache2
```

```
root@VM1-Server:~# sudo apt install vsftpd
```

```
root@VM1-Server:~# apt-get install openssh-server
```

Enabling opening ports on boot up:

```
root@VM1-Server:~# systemctl enable ssh
```

```
root@VM1-Server:~# systemctl enable apache2
```

```
root@VM1-Server:~# apt install vsftpd
```

```
root@VM1-Server:~# cp /etc/vsftpd.conf /etc/vsftpd.conf.bak
root@VM1-Server:~# nano /etc/vsftpd.conf
```

```
# Run standalone? vsftpd can run either from an inetd or as a standalone
# daemon started from an initscript.
listen=NO
#
# This directive enables listening on IPv6 sockets. By default, listening
# on the IPv6 "any" address (:::) will accept connections from both IPv6
# and IPv4 clients. It is not necessary to listen on *both* IPv4 and IPv6
# sockets. If you want that (perhaps because you want to listen on specific
# addresses) then you must run two copies of vsftpd with two configuration
# files.
listen_ipv6=YES
#
# Allow anonymous FTP? (Disabled by default).
anonymous_enable=YES
#
# Uncomment this to allow local users to log in.
local_enable=YES
```

### 2.3.3. VM2- Firewall settings

```
root@VM2-Firewall:~# ip a
1: lo: <LOOPBACK,UP,LOWER_UP> mtu 65536 qdisc noqueue state UNKNOWN group default qlen 1000
    link/loopback 00:00:00:00:00:00 brd 00:00:00:00:00:00
    inet 127.0.0.1/8 scope host lo
        valid_lft forever preferred_lft forever
    inet6 ::1/128 scope host
        valid_lft forever preferred_lft forever
2: enp0s3: <BROADCAST,MULTICAST,UP,LOWER_UP> mtu 1500 qdisc fq_codel state UP group default qlen 1000
    link/ether 08:00:27:6a:e7:c3 brd ff:ff:ff:ff:ff:ff
    inet 10.5.5.1/24 brd 10.5.5.255 scope global noprefixroute enp0s3
        valid_lft forever preferred_lft forever
    inet6 fe80::265d:7f91:d811:8667/64 scope link noprefixroute
        valid_lft forever preferred_lft forever
3: enp0s8: <BROADCAST,MULTICAST,UP,LOWER_UP> mtu 1500 qdisc fq_codel state UP group default qlen 1000
    link/ether 08:00:27:1c:c4:3a brd ff:ff:ff:ff:ff:ff
    inet 192.168.5.1/24 brd 192.168.5.255 scope global noprefixroute enp0s8
        valid_lft forever preferred_lft forever
    inet6 fe80::e021:3ee8:35dc:c95e/64 scope link noprefixroute
        valid_lft forever preferred_lft forever
```

Installing and enabling nftables on VM2Firewall and ssh, ftp services and port forwarding, examples below:

```
root@VM2-Firewall:~# sudo apt install nftables
root@VM2-Firewall:~# echo 1 > /proc/sys/net/ipv4/ip_forward
root@VM2-Firewall:~# systemctl enable ssh
root@VM2-Firewall:~# apt install vsftpd
Processing triggers for man-db (2.10.2-1) ...
root@VM2-Firewall:~# cp /etc/vsftpd.conf /etc/vsftpd.conf.bak
root@VM2-Firewall:~# nano /etc/vsftpd.conf
root@VM2-Firewall:~# systemctl start vsftpd
```

### 2.3.4. VM3 - Attacker settings

```
root@VM3-Attacker:~# ip a
1: lo: <LOOPBACK,UP,LOWER_UP> mtu 65536 qdisc noqueue state UNKNOWN group default qlen 1000
    link/loopback 00:00:00:00:00:00 brd 00:00:00:00:00:00
    inet 127.0.0.1/8 scope host lo
        valid_lft forever preferred_lft forever
    inet6 ::1/128 scope host
        valid_lft forever preferred_lft forever
2: enp0s3: <BROADCAST,MULTICAST,UP,LOWER_UP> mtu 1500 qdisc fq_codel state UP group default qlen 1000
    link/ether 08:00:27:8c:70:69 brd ff:ff:ff:ff:ff:ff
    inet 192.168.5.30/24 brd 192.168.5.255 scope global noprefixroute enp0s3
        valid_lft forever preferred_lft forever
    inet6 fe80::2484:df50:cdc5:583f/64 scope link noprefixroute
        valid_lft forever preferred_lft forever
```

Tools:

```
root@VM3-Attacker:~# sudo apt-get install nmap
root@VM3-Attacker:~# sudo apt install curl
root@VM3-Attacker:~# apt install net-tools
root@VM3-Attacker:~# sudo apt install hping3
```

Attacker can reach Firewall and Apache2 service in VM1-Server:

```
vm3attackeruser@VM3-Attacker:~$ ping 192.168.5.1
```

```

rtt min/avg/max/mdev = 1.223/1.611/2.154/0.277 ms
vm3attackeruser@VM3-Attacker:~$ curl 10.5.5.30
<!DOCTYPE html PUBLIC "-//W3C//DTD XHTML 1.0 Transitional//EN" "http://www.w3.org/
/xhtml1/DTD/xhtml1-transitional.dtd">
<html xmlns="http://www.w3.org/1999/xhtml">
  <!--
    Modified from the Debian original for Ubuntu
    Last updated: 2022-03-22
    See: https://launchpad.net/bugs/1966004
  -->
<head>

```

### 3. Testing

#### Summary of tests:

Author: Sebastian Konefal student no:b00168561		Tests of Nftable				
Number	Type	Tool	Source	Destination	Result	Comment
Test-3.1	Open ports scan	Nmap	VM3-Attacker	VM1-Server	PASS	NFTable blocked port scan discovery by frequent sources timeouts
	<i>nmap -p 1-65535 10.5.5.30</i>			NFTable ON		
Test-3.2	Open ports scan	Nmap	VM3-Attacker	VM1-Server	FAIL	NFTable was turned off and the scan revealed open services
	<i>nmap -p 1-65535 10.5.5.30</i>			NFTable OFF		
Test-3.3	Open ports scan:	Nmap	VM3-Attacker	VM1-Firewall	PASS	NFTable blocked port scan discovery by frequent sources timeouts
	<i>nmap -p 1-65535 192.168.5.1</i>			NFTable ON		
Test-3.4	Open ports scan:	Nmap	VM3-Attacker	VM1-Firewall	FAIL	NFTable was turned off and the scan revealed open services
	<i>nmap -p 1-65535 192.168.5.1</i>			NFTable OFF		
Test-3.5	IPv6 open port scan on	Nmap & ping	VM3-Attacker	VM-Firewall	PASS	No connection via IPv6
	<i>nmap -6 -p 1-65535 fe80::e021:3ee8:35dc:c95e/64</i>			NFTable ON		
	<i>ping fe80::e021:3ee8:35dc:c95e</i>					
Test-3.6	IPv6 open port scan on	Ping6	VM3-Attacker	VM1-Server	PASS	No connection via IPv6
	<i>ping6 fe80::43eb:1933:6699:4ab2%enp0s3</i>			NFTable ON		
Test-3.7	Accessing web server using IP: 192.168.5.30 & 192.168.5.190	Curl	VM3-Attacker	VM1-Server	PASS	Packets sent and reply received
	<i>curl 10.5.5.30</i>			NFTable ON		
Test-3.8	DOS:	Hping3	VM3-Attacker	VM1-Server - port 80	PASS	Attack for 60 seconds, 1,540,273 packets transmitted but only 122 processed (sent and received)
	<i>hping3 -d 120 -S -w 64 -p 80 --flood 10.5.5.30</i>			NFTable ON		

Test-3.9	DOS: <i>hping3 -d 120 -S -w 64 -p 80 --flood 10.5.5.30</i>	Hping3	VM3-Attacker	VM1-Server – port 80 NFTable OFF	FAIL	Attack for 60 seconds, 1,500,179 packet transmitted resulting in 794,546 packets received and total 1,377,327 packets processed
Test-3.10	DOS using IP of non-privileged range: <i>hping3 -d 120 -S -w 64 -p 22 --flood 192.168.5.1</i>	Hping3	VM3-Attacker	VM2-Firewall – intnet port 22 NFTable ON	PASS	Attack for 60 seconds, 331,053 packet dropped out of 331,054 sent. Allowed 1 packet.
Test-3.11	DOS using IP of privileged range - IP 192.168.5.30 used: <i>hping3 -d 120 -S -w 64 -p 22 --flood 192.168.5.1</i>	Hping3	VM3-Attacker	VM2-Firewall – extnet port 22 NFTable ON	FAIL – FIXED BY TEST 3.12	Attack for 60 seconds, 4,831,512 packets sent and 409,103 received resulting in total of 700,209 packets sent and received from VM2-Firewall
Test-3.12 FIX	DOS using IP of privileged range - IP 192.168.5.30 used: <i>hping3 -d 120 -S -w 64 -p 22 --flood 192.168.5.1</i>	Hping3	VM3-Attacker	VM2-Firewall – extnet port 22 NFTable ON	PASS	Additional rule was added to fix Test-3.11. Attack for 60 seconds, 332,650 packets transmitted and processed only 62 packets
Test-3.13	DOS using IP of privileged range - IP 192.168.5.30 used: <i>hping3 -d 120 -S -w 64 -p 22 --flood 192.168.5.1</i>	Hping3	VM3-Attacker	VM2-Firewall – extnet port NFTable OFF	FAIL	Attack for 60 seconds, 6,121,501 packets transmitted and 401,484 packets accepted by VM2-Firewall
Test-3.14	Triggering egress hook by accessing web server <i>curl 10.5.5.30</i>	Curl	VM2-Firewall	VM1-Server – port 80 NFTable ON	PASS	Egress hook triggered
Test-3.15	Connecting to FTP (hidden open port) & SSH <i>nc -w1 -vz 10.5.5.30 21</i> <i>nc -w1 -vz 10.5.5.30 22</i> <i>nc -w1 -vz 192.168.5.1 21</i> <i>nc -w1 -vz 192.168.5.1 22</i>	Netcat	VM2-Firewall	VM1-Server & VM2-Firewall – port 21 & 22 NFTable ON	PASS	NFTable blocked open FTP port and allowed SSH traffic

### 3.1. Open port scan on VM1-Server

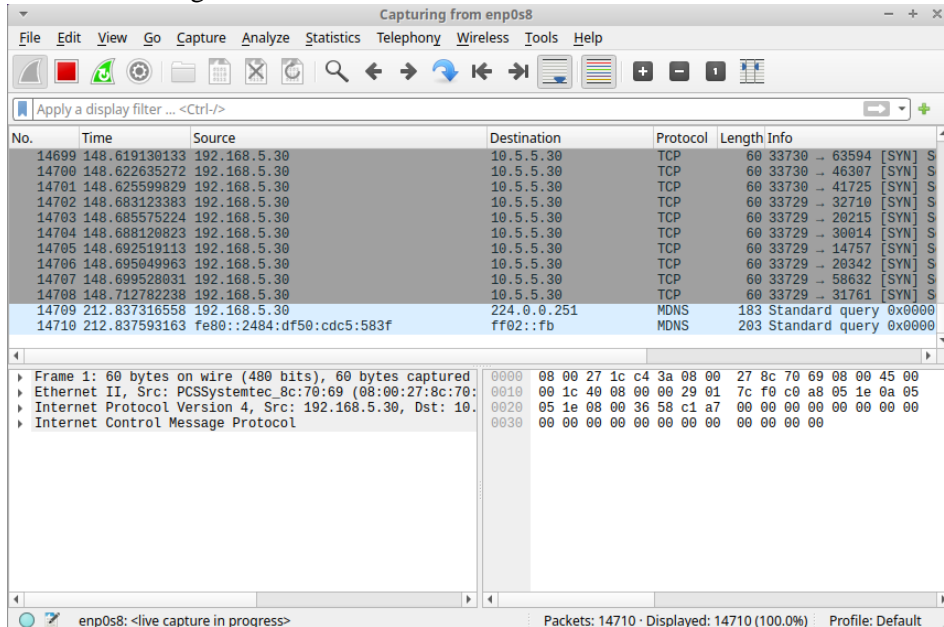
Nmap search for all possible ports. VM-Attacker does not see FTP open port in VM3-Firewall, because VM-Firewall is blocking discovery by timeouts:

```
root@VM3-Attacker:~# nmap -p 1-65535 10.5.5.30
Starting Nmap 7.80 ( https://nmap.org ) at 2024-03-03 15:54 GMT
root@VM3-Attacker:~# nmap -p 1-65535 192.168.5.1
```

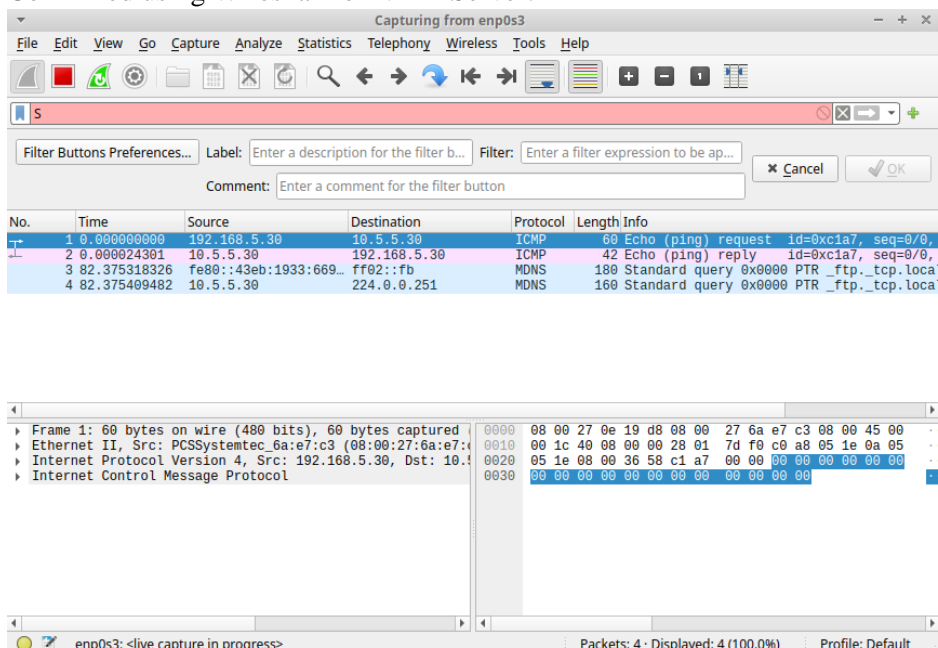
Confirmed block traffic in VM3-Firewall NFTable:

```
chain Serverforward {
    type filter hook forward priority filter; policy accept;
    ct state established,related accept
    ct state new ip saddr @frequent_server_sources counter packets 14543 bytes 639884 drop
    ct state new add @frequent_server_sources { ip saddr timeout 1s }
    ct state new tcp dport { 80, 443 } counter name "counter_ct_web" accept
    ip protocol icmp accept
    ct state invalid counter packets 0 bytes 0 drop
    tcp dport 22 jump input_ssh
    counter packets 147 bytes 6468 drop
}
```

Confirmed using Wireshark on VM2-Firewall:



Confirmed using Wireshark on VM1-Server:



**RESULT: PASS** - NfTable blocked port scan discovery by frequent sources timeouts

### 3.2. Open port scan on VM1-Server – **NFTable off**

```
root@VM3-Attacker:~# nmap -p 1-65535 10.5.5.30
Starting Nmap 7.80 ( https://nmap.org ) at 2024-03-03 16:11 GMT
Nmap scan report for 10.5.5.30
Host is up (0.00081s latency).
Not shown: 65532 closed ports
PORT      STATE SERVICE
21/tcp    open  ftp
22/tcp    open  ssh
80/tcp    open  http

Nmap done: 1 IP address (1 host up) scanned in 15.43 seconds
```

**RESULT: FAIL** – NFTable was turned off and the scan revealed additional open services

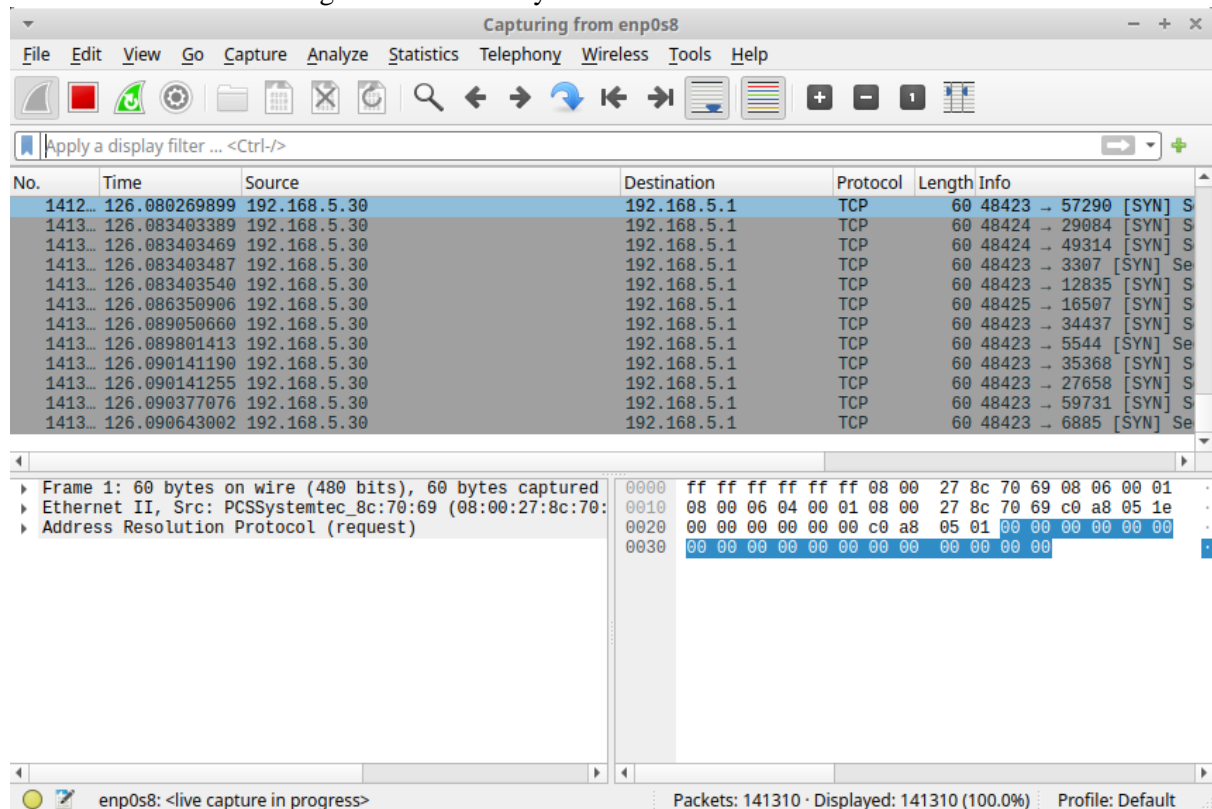
### 3.3. Open port scan on VM-Firewall

Nmap search for all possible ports. VM-Attacker does not see FTP open port in VM3-Firewall, because VM-Firewall is blocking discovery by timeouts:

```
MAC Address: 08:00:27:1C:C4:3A (Oracle VirtualBox virtual NIC)

Nmap done: 1 IP address (1 host up) scanned in 10.31 seconds
root@VM3-Attacker:~# nmap -p 1-65535 192.168.5.1
```

Confirmed traffic incoming VM2-Firewall by Wireshark:





Confirmed block traffic in VM3-Firewall NfTable:

```
chain input_firewall {
    type filter hook input priority filter; policy accept;
    iifname "lo" counter packets 0 bytes 0 accept
    ct state established,related counter packets 0 bytes 0 accept
    ct state new ip saddr @frequent_firewall_sources counter packets 140963 bytes 6202372 drop
    ct state new add @frequent_firewall_sources { ip saddr timeout 1s }
    ip protocol icmp accept
    ct state invalid counter packets 0 bytes 0 drop
    tcp dport 22 jump input_ssh
    counter packets 130 bytes 6124 drop
}
```

**RESULT: PASS** - NfTable blocked port scan discovery by frequent sources timeouts

### 3.4. Open port scan on VM2-Firewall – **NfTable off**

```
root@VM3-Attacker:~# nmap -p 1-65535 192.168.5.1
Starting Nmap 7.80 ( https://nmap.org ) at 2024-03-03 16:03 GMT
Nmap scan report for _gateway (192.168.5.1)
Host is up (0.00061s latency).
Not shown: 65533 closed ports
PORT      STATE SERVICE
21/tcp    open  ftp
22/tcp    open  ssh
MAC Address: 08:00:27:1C:C4:3A (Oracle VirtualBox virtual NIC)
Nmap done: 1 IP address (1 host up) scanned in 10.31 seconds
```

**RESULT: FAIL** – NfTable was turned off and the scan revealed additional open FTP service

### 3.5. IPv6 open port scan on VM-Firewall

```
root@VM3-Attacker:~# nmap -6 -p 1-65535 fe80::e021:3ee8:35dc:c95e/64
Starting Nmap 7.80 ( https://nmap.org ) at 2024-03-03 14:04 GMT
root@VM3-Attacker:~# ^C
```

No connection via ipv6:

```
root@VM3-Attacker:~# ping fe80::e021:3ee8:35dc:c95e
PING fe80::e021:3ee8:35dc:c95e(fe80::e021:3ee8:35dc:c95e) 56 data bytes
^C
--- fe80::e021:3ee8:35dc:c95e ping statistics ---
104 packets transmitted, 0 received, 100% packet loss, time 105451ms
```

## Packets transmitted via VM2-Firewall:

The screenshot shows a Wireshark capture on interface \*enp0s8. The display filter is empty. The packet list shows a series of ICMPv6 Neighbor Solicitation packets from source fe80::2484:df50:cdc5:583f to destination ff02::1:ff00:99db through ff02::1:ff00:9a0e. Packet 77983 is selected, showing a Router Solicitation of 62 bytes. The packet details pane shows the Ethernet II frame from PCSSystemtec\_1c:c4:3a (08:00:27:1c:c4:3a) to fe80::e021:3ee8:35dc:c95e, and the ICMPv6 Router Solicitation.

No.	Time	Source	Destination	Protocol	Length	Info
77977	1580.6233615...	fe80::2484:df50:cdc5:583f	ff02::1:ff00:99db	ICMPv6	86	Neighbor Solicitation
77978	1580.6267873...	fe80::2484:df50:cdc5:583f	ff02::1:ff00:99df	ICMPv6	86	Neighbor Solicitation
77979	1580.6297668...	fe80::2484:df50:cdc5:583f	ff02::1:ff00:99e4	ICMPv6	86	Neighbor Solicitation
77980	1580.6299639...	fe80::2484:df50:cdc5:583f	ff02::1:ff00:99e5	ICMPv6	86	Neighbor Solicitation
77981	1580.6301509...	fe80::2484:df50:cdc5:583f	ff02::1:ff00:99e8	ICMPv6	86	Neighbor Solicitation
77982	1580.6303659...	fe80::2484:df50:cdc5:583f	ff02::1:ff00:99e9	ICMPv6	86	Neighbor Solicitation
77983	1580.7758418...	fe80::e021:3ee8:35dc:c95e	ff02::2	ICMPv6	62	Router Solicitation
77984	1580.8180692...	fe80::2484:df50:cdc5:583f	ff02::1:ff00:9a02	ICMPv6	86	Neighbor Solicitation
77985	1580.8180696...	fe80::2484:df50:cdc5:583f	ff02::1:ff00:9a03	ICMPv6	86	Neighbor Solicitation
77986	1580.8214378...	fe80::2484:df50:cdc5:583f	ff02::1:ff00:9a07	ICMPv6	86	Neighbor Solicitation
77987	1580.8216422...	fe80::2484:df50:cdc5:583f	ff02::1:ff00:9a08	ICMPv6	86	Neighbor Solicitation
77988	1580.8249109...	fe80::2484:df50:cdc5:583f	ff02::1:ff00:9a0b	ICMPv6	86	Neighbor Solicitation
77989	1580.8274884...	fe80::2484:df50:cdc5:583f	ff02::1:ff00:9a0e	ICMPv6	86	Neighbor Solicitation

Frame 77983: 62 bytes on wire (496 bits), 62 bytes captured on interface \*enp0s8, 0 bytes filtered  
 Ethernet II, Src: PCSSystemtec\_1c:c4:3a (08:00:27:1c:c4:3a), Dst: fe80::e021:3ee8:35dc:c95e  
 Internet Protocol Version 6, Src: fe80::e021:3ee8:35dc:c95e, Dst: ff02::2  
 Internet Control Message Protocol v6

wireshark\_enp0s8KYZJ2.pcapng      Packets: 113811 · Displayed: 113811 (100.0%)      Profile: Default

## Packets replied by VM2-Firewall:

The screenshot shows the same Wireshark capture with a display filter applied: `ipv6.addr == fe80::e021:3ee8:35dc:c95e`. Only two packets are displayed: packet 52063 (MDNS Standard query) and packet 77983 (ICMPv6 Router Solicitation). The packet details pane for packet 77983 shows the Ethernet II frame from PCSSystemtec\_1c:c4:3a (08:00:27:1c:c4:3a) to fe80::e021:3ee8:35dc:c95e, and the ICMPv6 Router Solicitation.

No.	Time	Source	Destination	Protocol	Length	Info
52063	1057.1524765...	fe80::e021:3ee8:35dc:c95e	ff02::fb	MDNS	203	Standard query 0x0000
77983	1580.7758418...	fe80::e021:3ee8:35dc:c95e	ff02::2	ICMPv6	62	Router Solicitation

Frame 77983: 62 bytes on wire (496 bits), 62 bytes captured on interface \*enp0s8, 0 bytes filtered  
 Ethernet II, Src: PCSSystemtec\_1c:c4:3a (08:00:27:1c:c4:3a), Dst: fe80::e021:3ee8:35dc:c95e  
 Internet Protocol Version 6, Src: fe80::e021:3ee8:35dc:c95e, Dst: ff02::2  
 Internet Control Message Protocol v6

"fe80::e021:3ee8:35dc:c95e" is not a valid hostname or IPv6 address.      Packets: 113811 · Displayed: 2 (0.0%)      Profile: Default

**RESULT: PASS** - No connection via IPv6



### 3.6. IPv6 open port scan on VM2-Server

Open ports:

```
root@VM1-Server:~# netstat -tuln
Active Internet connections (only servers)
Proto Recv-Q Send-Q Local Address           Foreign Address         State
tcp        0      0 127.0.0.1:33060         0.0.0.0:*               LISTEN
tcp        0      0 127.0.0.1:631          0.0.0.0:*               LISTEN
tcp        0      0 0.0.0.0:22             0.0.0.0:*               LISTEN
tcp        0      0 127.0.0.1:3306         0.0.0.0:*               LISTEN
tcp        0      0 127.0.0.53:53          0.0.0.0:*               LISTEN
tcp6       0      0 :::80                  :::*                    LISTEN
tcp6       0      0 :::22                  :::*                    LISTEN
tcp6       0      0 :::21                  :::*                    LISTEN
tcp6       0      0 :::1:631               :::*                    LISTEN
udp        0      0 127.0.0.53:53          0.0.0.0:*               *
udp        0      0 0.0.0.0:631            0.0.0.0:*               *
udp        0      0 0.0.0.0:5353           0.0.0.0:*               *
udp        0      0 0.0.0.0:44340          0.0.0.0:*               *
udp6       0      0 :::50802               :::*                    *
udp6       0      0 :::5353                :::*                    *
```

No connection via ipv6:

```
root@VM3-Attacker:~# ping6 fe80::43eb:1933:6699:4ab2%enp0s3
PING fe80::43eb:1933:6699:4ab2%enp0s3(fe80::43eb:1933:6699:4ab2%enp0s3) 56 data bytes
From fe80::2484:df50:cdc5:583f%enp0s3 icmp_seq=1 Destination unreachable: Address unreachable
From fe80::2484:df50:cdc5:583f%enp0s3 icmp_seq=2 Destination unreachable: Address unreachable
From fe80::2484:df50:cdc5:583f%enp0s3 icmp_seq=3 Destination unreachable: Address unreachable
From fe80::2484:df50:cdc5:583f%enp0s3 icmp_seq=4 Destination unreachable: Address unreachable
ping6: sendmsg: No route to host
From fe80::2484:df50:cdc5:583f%enp0s3 icmp_seq=5 Destination unreachable: Address unreachable
From fe80::2484:df50:cdc5:583f%enp0s3 icmp_seq=6 Destination unreachable: Address unreachable
From fe80::2484:df50:cdc5:583f%enp0s3 icmp_seq=8 Destination unreachable: Address unreachable
ping6: sendmsg: No route to host
From fe80::2484:df50:cdc5:583f%enp0s3 icmp_seq=9 Destination unreachable: Address unreachable
From fe80::2484:df50:cdc5:583f%enp0s3 icmp_seq=10 Destination unreachable: Address unreachable
From fe80::2484:df50:cdc5:583f%enp0s3 icmp_seq=12 Destination unreachable: Address unreachable
ping6: sendmsg: No route to host
From fe80::2484:df50:cdc5:583f%enp0s3 icmp_seq=13 Destination unreachable: Address unreachable
From fe80::2484:df50:cdc5:583f%enp0s3 icmp_seq=14 Destination unreachable: Address unreachable
^C
--- fe80::43eb:1933:6699:4ab2%enp0s3 ping statistics ---
16 packets transmitted, 0 received, 100% packet loss, time 15357ms
```

**RESULT: PASS** - No connection via IPv6

### 3.7. Accessing Web Server on VM1-Server

Tested with VM3-Attacker IP 192.168.5.30 match by this rule:

```
chain prerouting server {
    type nat hook prerouting priority filter; policy accept;
    iifname "enp0s8" ip saddr 192.168.5.30 tcp dport { 80, 443 } counter packets 1 bytes 60 dnat to 10.5.5.30:80-443
}
```

And 192.168.5.190 matched by this rule in the forward hook:

```
ct state new tcp dport { 80, 443 } counter name "counter_ct_web" accept
```

Result for both IPs:

```
root@VM3-Attacker:~# curl 10.5.5.30
<!DOCTYPE html PUBLIC "-//W3C//DTD XHTML 1.0 Transitional//EN" "http://www.w3.org/TR/xhtml1/DTD/xhtml1-transitional.dtd">
<html xmlns="http://www.w3.org/1999/xhtml">
<!--
```

**RESULT: PASS** - Packet sent and reply received

### 3.8. DOS on VM2-Server port 80

VM3-Attacker sent 1,540,273 packets:

```
root@VM3-Attacker:~# hping3 -d 120 -S -w 64 -p 80 --flood 10.5.5.30
HPING 10.5.5.30 (enp0s3 10.5.5.30): S set, 40 headers + 120 data bytes
hping in flood mode, no replies will be shown
^C
--- 10.5.5.30 hping statistic ---
1540273 packets transmitted, 0 packets received, 100% packet loss
round-trip min/avg/max = 0.0/0.0/0.0 ms
```

VM2-Firewall blocked flooding attack by this rule of @frequent\_server\_sources:

```
chain PreroutingServer {
    type nat hook prerouting priority filter; policy accept;
    iifname "enp0s8" ip saddr 192.168.5.30 tcp dport { 80, 443 } counter packets 1540274 bytes 246443840 dnat to 10.5.5.30:80-443
}

chain Serverforward {
    type filter hook forward priority filter; policy accept;
    ct state established,related accept
    ct state new ip saddr @frequent_server_sources counter packets 1540213 bytes 246434080 drop
    ct state new add @frequent_server_sources { ip saddr timeout 1s }
    ct state new tcp dport { 80, 443 } counter name "counter_ct_web" accept
    ip protocol icmp accept
    ct state invalid counter packets 0 bytes 0 drop
    tcp dport 22 jump input_ssh
    counter packets 0 bytes 0 drop
}
```

VM1-Server processed (sent and received) **only 122 packets**:

Filter: ip.src == 192.168.5.30 or ip.dst == 192.168.5.30

No.	Time	Source	Destination	Protocol	Length	Info
1	0.000000000	192.168.5.30	10.5.5.30	TCP	174	2455 → 80 [SYN] Seq=0 Win=64 Len=120 [T
3	0.001211095	192.168.5.30	10.5.5.30	TCP	60	2455 → 80 [RST] Seq=1 Win=0 Len=0
4	1.018367145	192.168.5.30	10.5.5.30	TCP	174	27126 → 80 [SYN] Seq=0 Win=64 Len=120 [
6	1.019224515	192.168.5.30	10.5.5.30	TCP	60	27126 → 80 [RST] Seq=1 Win=0 Len=0
7	2.015601142	192.168.5.30	10.5.5.30	TCP	174	52616 → 80 [SYN] Seq=0 Win=64 Len=120 [
9	2.016524128	192.168.5.30	10.5.5.30	TCP	60	52616 → 80 [RST] Seq=1 Win=0 Len=0
10	3.015196722	192.168.5.30	10.5.5.30	TCP	174	12902 → 80 [SYN] Seq=0 Win=64 Len=120 [
12	3.016137753	192.168.5.30	10.5.5.30	TCP	60	12902 → 80 [RST] Seq=1 Win=0 Len=0
13	4.045924049	192.168.5.30	10.5.5.30	TCP	174	38986 → 80 [SYN] Seq=0 Win=64 Len=120 [
15	4.046816971	192.168.5.30	10.5.5.30	TCP	60	38986 → 80 [RST] Seq=1 Win=0 Len=0
18	5.042701778	192.168.5.30	10.5.5.30	TCP	174	64495 → 80 [SYN] Seq=0 Win=64 Len=120 [
20	5.043500018	192.168.5.30	10.5.5.30	TCP	60	64495 → 80 [RST] Seq=1 Win=0 Len=0
23	6.042043830	192.168.5.30	10.5.5.30	TCP	174	24655 → 80 [SYN] Seq=0 Win=64 Len=120 [

Frame 1: 174 bytes on wire (1392 bits), 174 bytes captured  
Ethernet II, Src: PCSSystemtec\_6a:e7:c3 (08:00:27:6a:e7:c3), Dst: 10.5.5.30  
Internet Protocol Version 4, Src: 192.168.5.30, Dst: 10.5.5.30  
Transmission Control Protocol, Src Port: 2455, Dst Port: 80

enp0s3: <live capture in progress> Packets: 195 · Displayed: 122 (62.6%) Profile: Default

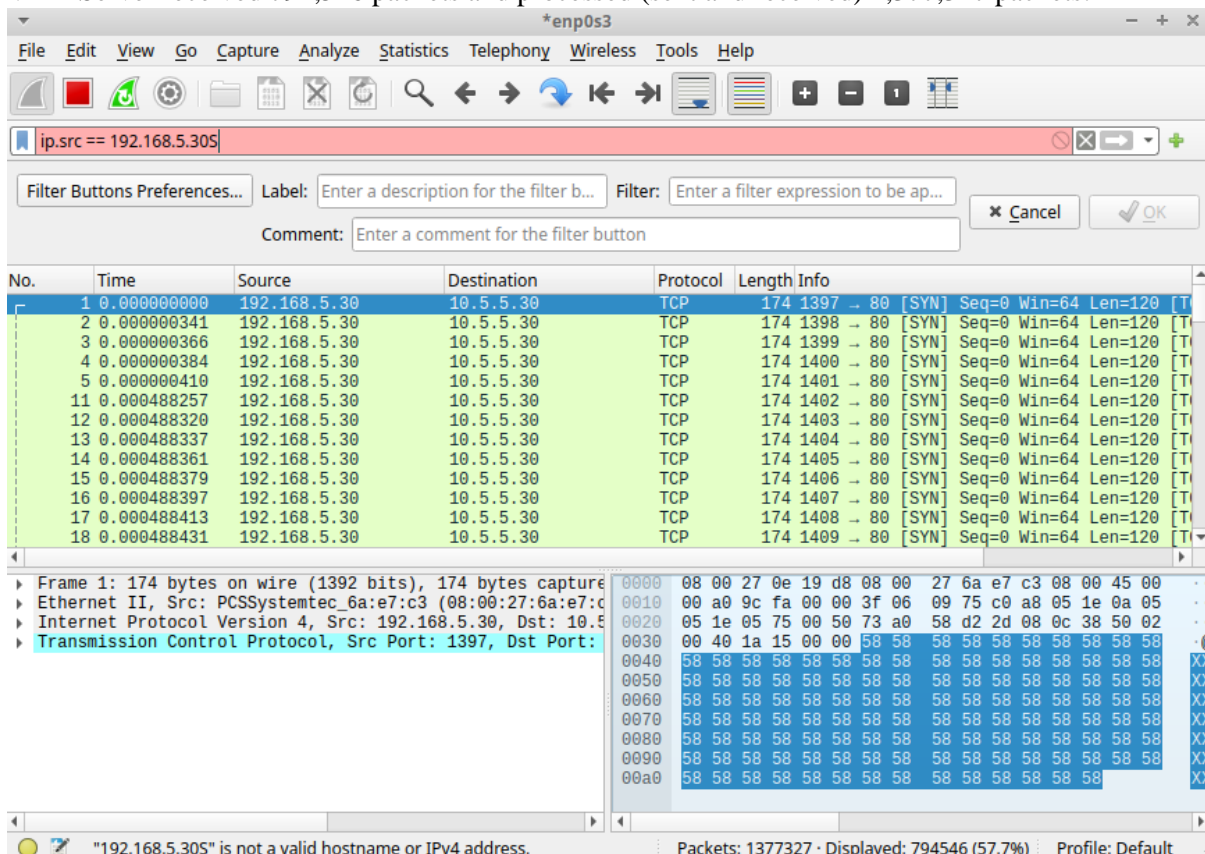
**RESULT: PASS** - Transmitted 1,540,273 packets but only 122 processed (sent and received)

### 3.9. DOS on VM2-Server port 80 – **NFTable off**

VM3-Attacker sent 1,500,179 packets:

```
root@VM3-Attacker:~# hping3 -d 120 -S -w 64 -p 80 --flood 10.5.5.30
HPING 10.5.5.30 (enp0s3 10.5.5.30): S set, 40 headers + 120 data bytes
hping in flood mode, no replies will be shown
^C
--- 10.5.5.30 hping statistic ---
1500179 packets transmitted, 0 packets received, 100% packet loss
round-trip min/avg/max = 0.0/0.0/0.0 ms
```

VM1-Server received 794,546 packets and processed (sent and received) 1,377,327 packets:



Filter: ip.src == 192.168.5.30

Filter Buttons Preferences... Label: Enter a description for the filter b... Filter: Enter a filter expression to be ap... Cancel OK

Comment: Enter a comment for the filter button

No.	Time	Source	Destination	Protocol	Length	Info
1	0.000000000	192.168.5.30	10.5.5.30	TCP	174	1397 → 80 [SYN] Seq=0 Win=64 Len=120
2	0.000000341	192.168.5.30	10.5.5.30	TCP	174	1398 → 80 [SYN] Seq=0 Win=64 Len=120
3	0.000000366	192.168.5.30	10.5.5.30	TCP	174	1399 → 80 [SYN] Seq=0 Win=64 Len=120
4	0.000000384	192.168.5.30	10.5.5.30	TCP	174	1400 → 80 [SYN] Seq=0 Win=64 Len=120
5	0.000000410	192.168.5.30	10.5.5.30	TCP	174	1401 → 80 [SYN] Seq=0 Win=64 Len=120
11	0.000488257	192.168.5.30	10.5.5.30	TCP	174	1402 → 80 [SYN] Seq=0 Win=64 Len=120
12	0.000488320	192.168.5.30	10.5.5.30	TCP	174	1403 → 80 [SYN] Seq=0 Win=64 Len=120
13	0.000488337	192.168.5.30	10.5.5.30	TCP	174	1404 → 80 [SYN] Seq=0 Win=64 Len=120
14	0.000488361	192.168.5.30	10.5.5.30	TCP	174	1405 → 80 [SYN] Seq=0 Win=64 Len=120
15	0.000488379	192.168.5.30	10.5.5.30	TCP	174	1406 → 80 [SYN] Seq=0 Win=64 Len=120
16	0.000488397	192.168.5.30	10.5.5.30	TCP	174	1407 → 80 [SYN] Seq=0 Win=64 Len=120
17	0.000488413	192.168.5.30	10.5.5.30	TCP	174	1408 → 80 [SYN] Seq=0 Win=64 Len=120
18	0.000488431	192.168.5.30	10.5.5.30	TCP	174	1409 → 80 [SYN] Seq=0 Win=64 Len=120

Frame 1: 174 bytes on wire (1392 bits), 174 bytes captured on interface enp0s3, 174 bytes from 192.168.5.30 to 10.5.5.30  
Ethernet II, Src: PCSSystemtec\_6a:e7:c3 (08:00:27:6a:e7:c3), Dst: 08:00:27:6a:e7:c3 (08:00:27:6a:e7:c3)  
Internet Protocol Version 4, Src: 192.168.5.30, Dst: 10.5.5.30  
Transmission Control Protocol, Src Port: 1397, Dst Port: 80

0000 08 00 27 0e 19 d8 08 00 27 6a e7 c3 08 00 45 00  
0010 00 a0 9c fa 00 00 3f 06 09 75 c0 a8 05 1e 0a 05  
0020 05 1e 05 75 00 50 73 a0 58 d2 2d 08 0c 38 50 02  
0030 00 40 1a 15 00 00 58 58 58 58 58 58 58 58 58  
0040 58 58 58 58 58 58 58 58 58 58 58 58 58 58 58  
0050 58 58 58 58 58 58 58 58 58 58 58 58 58 58 58  
0060 58 58 58 58 58 58 58 58 58 58 58 58 58 58 58  
0070 58 58 58 58 58 58 58 58 58 58 58 58 58 58 58  
0080 58 58 58 58 58 58 58 58 58 58 58 58 58 58 58  
0090 58 58 58 58 58 58 58 58 58 58 58 58 58 58 58  
00a0 58 58 58 58 58 58 58 58 58 58 58 58 58 58 58

"192.168.5.305" is not a valid hostname or IPv4 address. Packets: 1377327 · Displayed: 794546 (57.7%) Profile: Default

**RESULT: FAIL** - 1,500,179 packet transmitted resulting in 794,546 packets received and total 1,377,327 packets processed.

### 3.10. DOS on VM2-Firewall intnet on port 22 using IP of non-privileged range

Privileged range is:

```
#Decoupling list of IPs allowed to connect vis SSH
set allowed_ssh_ips{
    typeof ip saddr . tcp dport
    flags interval,constant #to use CIDR range and prevent changes from cl
    auto-merge #merge any overlapping range
    elements = {192.168.5.30/30 . 22}
}
```

I changed IP on VM3-Attacker to range outside of privileged range:

Address	Netmask	Gateway
192.168.5.190	24	192.168.5.1

Verified that connection is possible from VM3-Attacker but blocked on 3<sup>rd</sup> connection:

```
root@VM3-Attacker:~# nc -w1 -vz 192.168.5.1 22
Connection to 192.168.5.1 22 port [tcp/ssh] succeeded!
root@VM3-Attacker:~# nc -w1 -vz 192.168.5.1 22
Connection to 192.168.5.1 22 port [tcp/ssh] succeeded!
root@VM3-Attacker:~# nc -w1 -vz 192.168.5.1 22
^[[A
nc: connect to 192.168.5.1 port 22 (tcp) timed out: Operation now in progress
```

I reset the timeouts and VM3-Attacker sent 331,054 packets:

```
root@VM3-Attacker:~# hping3 -d 120 -S -w 64 -p 22 --flood 192.168.5.1
HPING 192.168.5.1 (enp0s3 192.168.5.1): S set, 40 headers + 120 data bytes
hping in flood mode, no replies will be shown
^C
--- 192.168.5.1 hping statistic ---
331054 packets transmitted, 0 packets received, 100% packet loss
round-trip min/avg/max = 0.0/0.0/0.0 ms
```

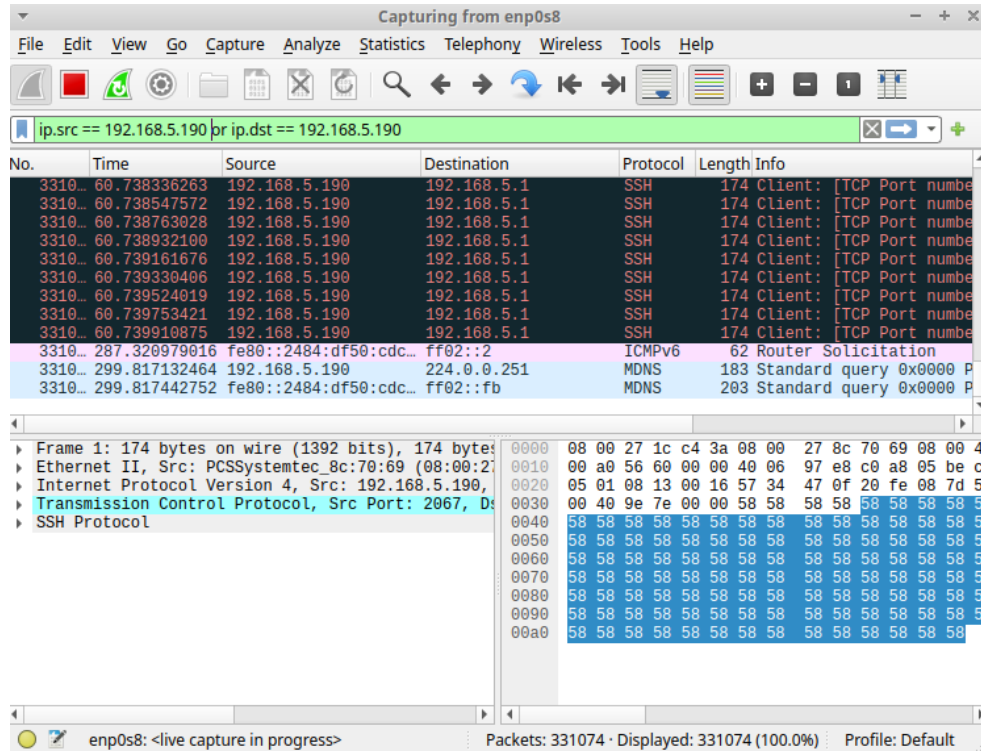
Received on VM2-Firewall and blocked by rule:

```
set timeout3 {
    typeof ip saddr
    size 65535
    flags dynamic,timeout
    elements = { 192.168.5.190 timeout 3d expires 2d23h53m6s508ms }
}
```

```
ct state new ip saddr . tcp dport @allowed_ssh_ips accept
ct state new ip saddr @timeout2 tcp dport 22 add @timeout3 { ip saddr timeout 3d }
ct state new ip saddr @timeout1 tcp dport 22 add @timeout2 { ip saddr timeout 3m }
ct state new tcp dport 22 add @timeout1 { ip saddr timeout 1m }
ct state new ip saddr @timeout3 tcp dport 22 counter packets 331053 bytes 52968480 drop
ct state new tcp dport 22 counter name "counter_ct_ssh" accept
counter packets 3 bytes 507 drop
```



Wireshark on VM2-Firewall:



**RESULT: PASS** - 331053 packet dropped out of 331054 send by VM3-Attacker. Allowed 1 packet.

### 3.11. DOS on VM2-Firewall on port 22 using IP of privileged range - IP 192.168.5.30 used

Privileged range is:

```
#Decoupling list of IPs allowed to connect vis SSH
set allowed_ssh_ips{
    type of ip saddr . tcp dport
    flags interval,constant #to use CIDR range and prevent changes from cl
    auto-merge #merge any overlapping range
    elements = {192.168.5.30/30 . 22}
}

#Dedicated rule to managment access via allowd range of IPs
ct state new ip saddr . tcp dport @allowed_ssh_ips accept #to prevent blocking management access
```

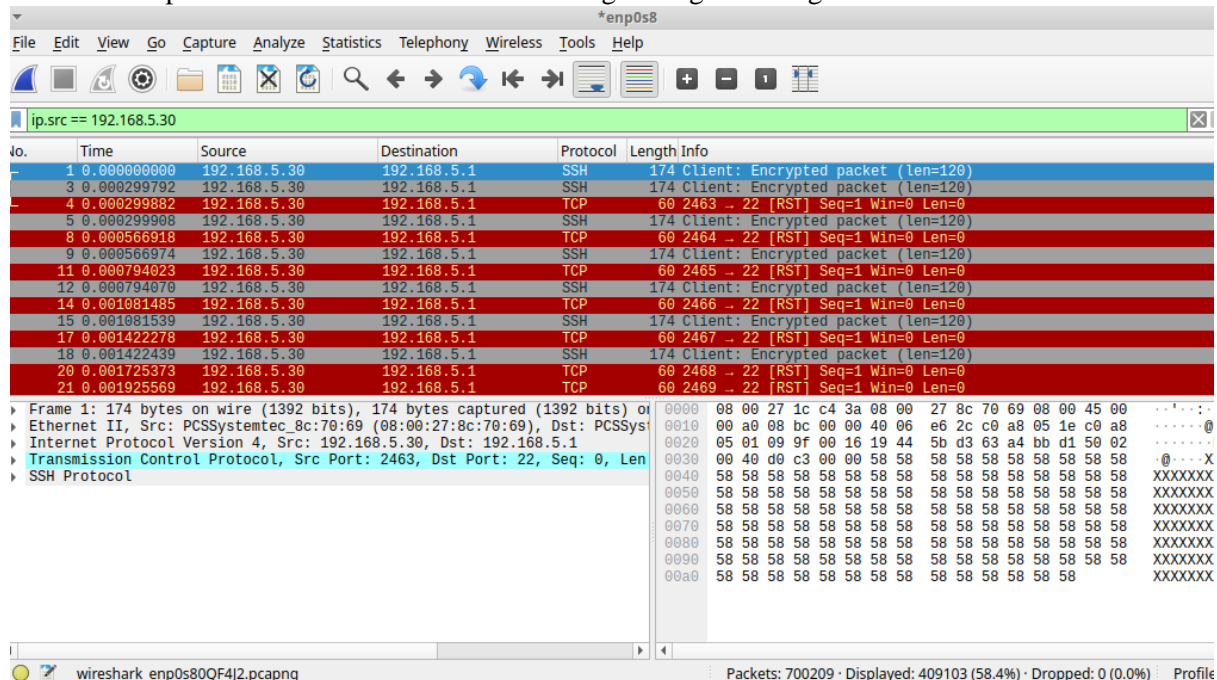
Flood attack from VM3-Attacker sending 4,831,512packets:

```
root@VM3-Attacker:~# hping3 -d 120 -S -w 64 -p 22 --flood 192.168.5.1
HPING 192.168.5.1 (enp0s3 192.168.5.1): S set, 40 headers + 120 data bytes
hping in flood mode, no replies will be shown
^C
--- 192.168.5.1 hping statistic ---
4831512 packets transmitted, 0 packets received, 100% packet loss
round-trip min/avg/max = 0.0/0.0/0.0 ms
```

There is an exception of admitted IP range on Nftable, therefore the entire traffic was allowed as it came from within such range:

```
ct state established,related counter packets 228470 bytes 25826360 accept
```

Wireshark capture on VM-Firewall 192.168.5.1 registering incoming traffic:



**RESULT: FAIL** 4,831,000 packets sent and 409,103 received resulting in total of 700,209 packets sent and received by VM2-Firewall

### 3.12. FIXED - DOS on VM2-Firewall on port 22 using IP of privileged range - IP 192.168.5.30 used

Test 3.12 was re-run with additional rule on input hook in VM2-Firewall:

```
#variable to prevent flood attack on firewall
set frequent_firewall_sources {
    type of ip saddr
    flags timeout
}
```

```
#Limiting flood attacks with 1s timeout
ct state new ip saddr @frequent_firewall_sources counter drop
ct state new add @frequent_firewall_sources {ip saddr timeout 1s}
```

Verified that connection is possible from VM3-Attacker but blocked on 3<sup>rd</sup> connection:

```
root@VM3-Attacker:~# nc -w1 -vz 192.168.5.1 22
Connection to 192.168.5.1 22 port [tcp/ssh] succeeded!
root@VM3-Attacker:~# nc -w1 -vz 192.168.5.1 22
Connection to 192.168.5.1 22 port [tcp/ssh] succeeded!
root@VM3-Attacker:~# nc -w1 -vz 192.168.5.1 22
nc: connect to 192.168.5.1 port 22 (tcp) timed out: Operation now in progress
```

I reset the timeouts. Flood attack from VM3-Attacker sending 332,650 packets:

```

root@VM3-Attacker:~# hping3 -d 120 -S -w 64 -p 22 --flood 192.168.5.1
HPING 192.168.5.1 (enp0s3 192.168.5.1): S set, 40 headers + 120 data bytes
hping in flood mode, no replies will be shown
^C
--- 192.168.5.1 hping statistic ---
332650 packets transmitted, 0 packets received, 100% packet loss
round-trip min/avg/max = 0.0/0.0/0.0 ms

```

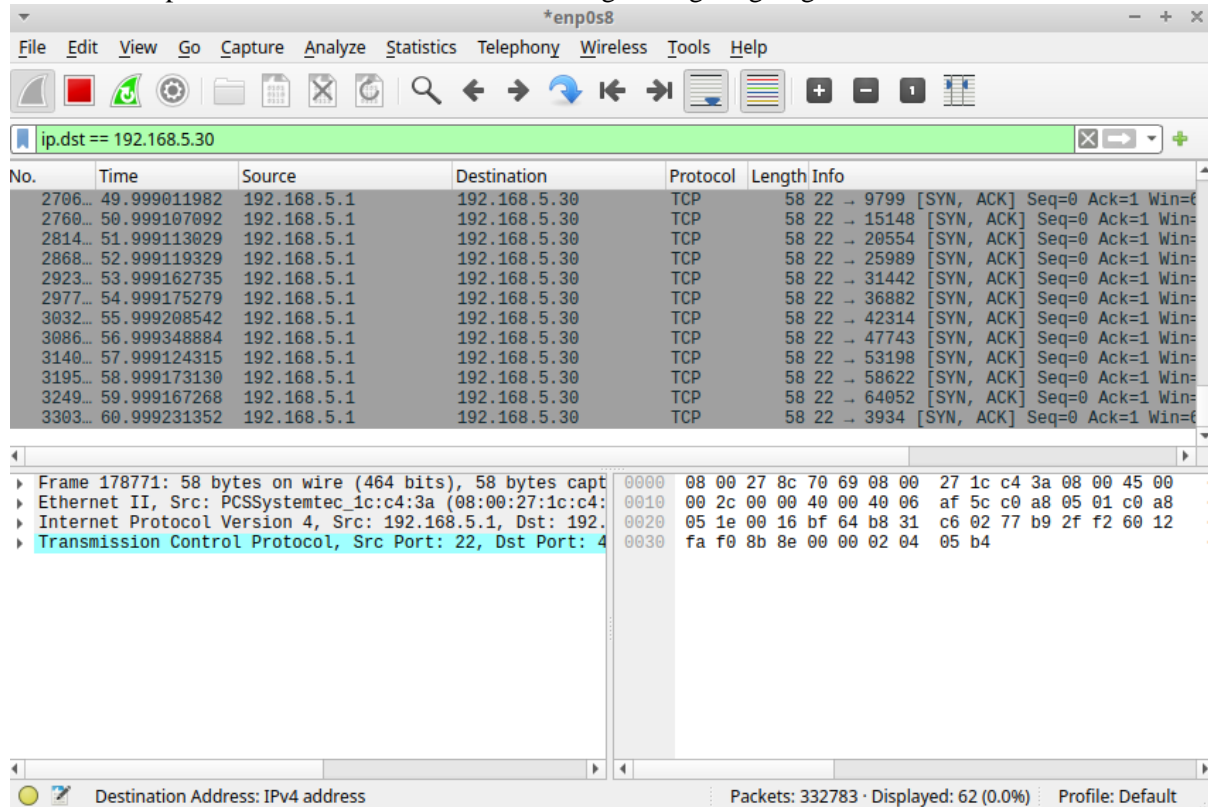
Wireshark capture on VM-Firewall 192.168.5.1 **registering incoming traffic:**

The Wireshark capture shows incoming traffic on interface \*enp0s8. The packet list displays a series of SSH connections from source 192.168.5.30 to destination 192.168.5.1. The packet details pane for the selected packet (Frame 332775) shows the following structure:

- Frame 332775: 174 bytes on wire (1392 bits), 174 bytes captured (1392 bits) on interface \*enp0s8
- Ethernet II, Src: PCSSystemtec\_8c:70:69 (08:00:27:8c:70:69), Dst: 192.168.5.1 (08:00:27:8c:70:69)
- Internet Protocol Version 4, Src: 192.168.5.30, Dst: 192.168.5.1
- Transmission Control Protocol, Src Port: 6332, Dst Port: 22
- SSH Protocol

The packet bytes pane shows the raw data of the selected packet, starting with the Ethernet II header (08 00 27 1c c4 3a 08 00 27 8c 70 69 08 00 45 00) and the IP header (00 10 00 a0 bf 7a 00 00 40 06 2f 6e c0 a8 05 1e c0 a8).

Wireshark capture on VM-Firewall 192.168.5.1 registering outgoing traffic:



Received on VM2-Firewall and blocked by rule:

```
ct state established,related counter packets 62 bytes 2480 accept
ct state new ip saddr @frequent_firewall_sources counter packets 332589 bytes 53214240 drop
```

**RESULT: PASS** - 332650 packets transmitted by VM3-Attacker and only 62 packets processed by VM2-Firewall

### 3.13. DOS on VM2-Firewall on port 22 using IP of privileged range - IP 192.168.5.30 used with Nftable off

Flood attack from VM3-Attacker sending 6,121,501 packets:

```
root@VM3-Attacker:~# hping3 -d 120 -S -w 64 -p 22 --flood 192.168.5.1
hping3: you must specify only one target host at a time
root@VM3-Attacker:~# hping3 -d 120 -S -w 64 -p 22 --flood 192.168.5.1
HPING 192.168.5.1 (enp0s3 192.168.5.1): S set, 40 headers + 120 data bytes
hping in flood mode, no replies will be shown
^C
--- 192.168.5.1 hping statistic ---
6121501 packets transmitted, 0 packets received, 100% packet loss
round-trip min/avg/max = 0.0/0.0/0.0 ms
```



### Wireshark capture on VM-Firewall 192.168.5.1 registering incoming traffic:

Wireshark capture on VM-Firewall 192.168.5.1 showing incoming traffic. The filter is `ip.src == 192.168.5.30`. The packet list shows 19 packets, including SSH and TCP connections. The packet details pane shows the structure of the first packet: Ethernet II, Internet Protocol Version 4, and Transmission Control Protocol.

No.	Time	Source	Destination	Protocol	Length	Info
1	0.000000000	192.168.5.30	192.168.5.1	SSH	174	Client: Encrypted pac
3	0.000305422	192.168.5.30	192.168.5.1	SSH	174	Client: Encrypted pac
5	0.000916667	192.168.5.30	192.168.5.1	TCP	60	2010 → 22 [RST] Seq=1
6	0.000916857	192.168.5.30	192.168.5.1	TCP	60	2011 → 22 [RST] Seq=1
7	0.000916905	192.168.5.30	192.168.5.1	SSH	174	Client: Encrypted pac
9	0.001232324	192.168.5.30	192.168.5.1	SSH	174	Client: Encrypted pac
10	0.001232391	192.168.5.30	192.168.5.1	TCP	60	2012 → 22 [RST] Seq=1
11	0.001232410	192.168.5.30	192.168.5.1	SSH	174	Client: Encrypted pac
14	0.001671302	192.168.5.30	192.168.5.1	TCP	60	2013 → 22 [RST] Seq=1
15	0.001671359	192.168.5.30	192.168.5.1	TCP	60	2014 → 22 [RST] Seq=1
16	0.001671399	192.168.5.30	192.168.5.1	SSH	174	Client: Encrypted pac
18	0.002015010	192.168.5.30	192.168.5.1	SSH	174	Client: Encrypted pac
19	0.002015073	192.168.5.30	192.168.5.1	TCP	60	2015 → 22 [RST] Seq=1

Frame 1: 174 bytes on wire (1392 bits), 174 bytes captured  
 Ethernet II, Src: PCSSystemtec\_8c:70:69 (08:00:27:8c:70:69), Dst: 08:00:27:8c:70:69  
 Internet Protocol Version 4, Src: 192.168.5.30, Dst: 192.168.5.1  
 Transmission Control Protocol, Src Port: 2010, Dst Port: 22  
 SSH Protocol

"S" was unexpected in this context. Packets: 711871 · Displayed: 401484 (56.4%) Profile: Default

### Wireshark capture on VM-Firewall 192.168.5.1 registering outgoing traffic:

Wireshark capture on VM-Firewall 192.168.5.1 showing outgoing traffic. The filter is `ip.dst == 192.168.5.30`. The packet list shows 17 packets, all TCP connections. The packet details pane shows the structure of the second packet: Ethernet II, Internet Protocol Version 4, and Transmission Control Protocol.

No.	Time	Source	Destination	Protocol	Length	Info
2	0.00044921	192.168.5.1	192.168.5.30	TCP	58	22 → 2010 [SYN, ACK]
4	0.000341591	192.168.5.1	192.168.5.30	TCP	58	22 → 2011 [SYN, ACK]
8	0.000929288	192.168.5.1	192.168.5.30	TCP	58	22 → 2012 [SYN, ACK]
12	0.001239229	192.168.5.1	192.168.5.30	TCP	58	22 → 2013 [SYN, ACK]
13	0.001274147	192.168.5.1	192.168.5.30	TCP	58	22 → 2014 [SYN, ACK]
17	0.001678107	192.168.5.1	192.168.5.30	TCP	58	22 → 2015 [SYN, ACK]
20	0.002020585	192.168.5.1	192.168.5.30	TCP	58	22 → 2016 [SYN, ACK]
24	0.002408237	192.168.5.1	192.168.5.30	TCP	58	22 → 2017 [SYN, ACK]
25	0.002449725	192.168.5.1	192.168.5.30	TCP	58	22 → 2018 [SYN, ACK]
30	0.002777218	192.168.5.1	192.168.5.30	TCP	58	22 → 2019 [SYN, ACK]
31	0.002874276	192.168.5.1	192.168.5.30	TCP	58	22 → 2020 [SYN, ACK]
35	0.003206034	192.168.5.1	192.168.5.30	TCP	58	22 → 2021 [SYN, ACK]
38	0.003427624	192.168.5.1	192.168.5.30	TCP	58	22 → 2022 [SYN, ACK]

Frame 2: 58 bytes on wire (464 bits), 58 bytes captured  
 Ethernet II, Src: PCSSystemtec\_1c:c4:3a (08:00:27:1c:c4:3a), Dst: 08:00:27:8c:70:69  
 Internet Protocol Version 4, Src: 192.168.5.1, Dst: 192.168.5.30  
 Transmission Control Protocol, Src Port: 22, Dst Port: 2010

Destination Address: IPv4 address Packets: 711874 · Displayed: 310375 (43.6%) Profile: Default

**RESULT: FAIL** - 6,121,501 packets transmitted by VM3-Attacker and 401,484 packets accepted by VM2-Firewall

### 3.14. Triggering egress hook by accessing web server

```
root@VM2-Firewall:~# curl 10.5.5.30
<!DOCTYPE html PUBLIC "-//W3C//DTD XHTML 1.0 Transitional//EN" "http://www.w3.org/TR/xhtml1/DTD/xhtml1-transitional.dtd">
<html xmlns="http://www.w3.org/1999/xhtml">
<!--
  Modified from the Debian original for Ubuntu
  Last updated: 2022-03-22
  See: https://launchpad.net/bugs/1966004
-->
<head>
```

Triggered egress hook:

```
set egress_tcp_connections {
    typeof meta skuid . ip daddr . tcp dport
    size 65535
    flags dynamic
    counter
    elements = { 0 . 10.5.5.30 . 80 counter packets 1 bytes 60 }
}
```

```
chain egress {
    type filter hook output priority filter; policy accept;
    ct state new add @egress_tcp_connections { meta skuid . ip daddr . tcp dport }
}
```

**RESULT: PASS** – Egress hook triggered

### 3.15. Connecting to FTP (hidden open port) & SSH

Both FTP and SSH ports are open as discovered by nmap in test 3.2 & 3.4 when NFtable was off.

Tests:

```
root@VM3-Attacker:~# nc -w1 -vz 10.5.5.30 21
nc: connect to 10.5.5.30 port 21 (tcp) timed out: Operation now in progress
```

```
root@VM3-Attacker:~# nc -w1 -vz 10.5.5.30 22
Connection to 10.5.5.30 22 port [tcp/ssh] succeeded!
```

```
root@VM3-Attacker:~# nc -w1 -vz 192.168.5.1 21
nc: connect to 192.168.5.1 port 21 (tcp) timed out: Operation now in progress
```

```
root@VM3-Attacker:~# nc -w1 -vz 192.168.5.1 22
Connection to 192.168.5.1 22 port [tcp/ssh] succeeded!
```

**RESULT: PASS** NFtable blocked open FTP port and allowed SSH traffic

## 4. Conclusions

The assignment involved developing a testbed scenario involving 3VMs with Linux OS simulating an attacker, and a firewall in between the communication and server.

In the first stage, I had to configure all VMs to enable FTP, SSH and Web services and communication through the VM2-Fierwall.

Subsequently, I learned about and configured NFtables to establish a robust firewall, implemented security measures, and finally conducted various tests. I developed key components such as port filtering, SSH & Web access control, a prerouting hook for Destination Network Address Translation (DNAT), and an egress rule.

In my project, I designed the ruleset to prevent flooding attacks, managed timeouts for SSH connections, defined privileged SSH access for management, and developed a forwarding hook for specific incoming traffic to a designated server.

In the testing stage, I covered scenarios such as port scans, SSH flood attacks depending on source IP, Web flood attacks, IPv6 accessibility, connecting via Netcat and egress hook triggering.

Continuous testing and refinement are crucial to adapting to evolving threats and ensuring the effectiveness of the firewall configuration. As I mentioned before in the report, further development might consist of developing Honeypots, Logging and Analysis, Geo-IP Filtering, Load Balancing.

Certainly, the most time-consuming was developing firewall logic and troubleshooting/debugging including in the testing stage.

## 5. Source editable NFTable:

GNU nano 6.2

firewall.nft

flush ruleset

table ip Server{

#Variables:

#Decoupling list of IPs allowed to connect via SSH

set allowed\_ssh\_ips{

    typeof ip saddr . tcp dport

    flags interval,constant #to use CIDR range and prevent changes from cl

    auto-merge #merge any overlapping range

    elements = { 192.168.5.30/30 . 22 }

}

#Counter variables

counter counter\_ct\_web{

}

counter counter\_ct\_ssh{

}

#3x timeout variables

set timeout1{

    typeof ip saddr

    flags timeout

}

set timeout2{

    typeof ip saddr

    flags timeout

}

set timeout3{

    typeof ip saddr

    flags timeout

```

}

#Allowed server ports
set allowed_server_ports{
    typeof tcp dport
    elements = { 22, 80, 443 }
}

#Variable to prevent flood attack on server
set frequent_server_sources {
    typeof ip saddr
    flags timeout
}

#variable to prevent flood attack on firewall
set frequent_firewall_sources {
    typeof ip saddr
    flags timeout
}

#CHAINS:

chain PreroutingServer{
    type nat hook prerouting priority 0;
    #Rule to forward only specified IP - implemented for testing purposes
    iifname "enp0s8" ip saddr 192.168.5.30 tcp dport { 80, 443 } counter dnat to 10.5.5.30:80-
443
    #continue to statefull packet examination
}

chain Serverforward {
    type filter hook forward priority 0;
    # Permit established and related SSH connections

```

```

ct state established, related accept

#Limiting new flood attack with 1s timeout
ct state new ip saddr @frequent_server_sources counter drop
ct state new add @frequent_server_sources {ip saddr timeout 1s}
#Forawrding web services
ct state new tcp dport {80, 443 }counter name counter_ct_web accept
#Accepting icmp
ip protocol icmp accept
#Dropping malicious/invalid packets
ct state invalid counter drop

#JUMP to chain dealing with SSH port 22
tcp dport 22 jump input_ssh

#Dropping any other traffic
counter drop
}

chain input_firewall {
    type filter hook input priority filter; policy accept;
    iifname lo counter accept
    #if a contract receives a SYN packet it considers connection as new , SYN-ACK considers
as established
    ct state established, related counter accept

    #Limiting flood attacks with 1s timeout
    ct state new ip saddr @frequent_firewall_sources counter drop
    ct state new add @frequent_firewall_sources {ip saddr timeout 1s}

    #Accepting icmp
    ip protocol icmp accept
    #Dropping malicious/invalid packets

```

```

ct state invalid counter drop

#JUMP to chain dealing with SSH port 22
tcp dport 22 jump input_ssh

#Dropping any other traffic
counter drop
}

chain input_ssh{
    # Permit established and related SSH connections
    ct state established, related accept

    #Dedicated rule to managment access via allowed range of IPs -it will accept unlimited
connections but not more often

    #then every 1s as per limiting flood attack rule in input_firewall chain
    ct state new ip saddr . tcp dport @allowed_ssh_ips counter accept #to prevent blocking
management access

    #Time-outs for SSH
    ct state new ip saddr @timeout2 tcp dport 22 add @timeout3 {ip saddr timeout 3d}
    ct state new ip saddr @timeout1 tcp dport 22 add @timeout2 {ip saddr timeout 3m}
    ct state new tcp dport 22 add @timeout1 {ip saddr timeout 1m}
    ct state new ip saddr @timeout3 tcp dport 22 counter drop

    #Only SYN packet will match this rule
    ct state new tcp dport 22 counter name counter_ct_ssh accept

}

```

## #OUTPUT

```

#variable to track outgoing connections:
set egress_tcp_connections{
    typeof meta skuid . ip daddr . tcp dport;
    counter;
}

chain egress {

```

```
type filter hook output priority filter; policy accept;
ct state new ip protocol tcp add @egress_tcp_connections {meta skuid . ip daddr . tcp dport}
}

}
```



## 6. References:

1. Suehring S., (2015) Linux firewalls: enhancing security with nftables and beyond 4<sup>th</sup> edition, Addison-Wesley
2. Shirvar A., (2020) ‘A comprehensive guide to Nftables (A leading packet filtering framework for Linux)’ in linkedin.com [online]. Available at: <https://www.linkedin.com/pulse/comprehensive-guide-nftables-leading-packet-filtering-arash-shirvar> [accessed 26 February 2024]
3. SW Team, (2024) ‘Nftables: a deep look at linux security’ in swhosting.com [online]. Available at: <https://www.swhosting.com/en/blog/nftables-a-deep-look-at-linux-security> [accessed 26 February 2024]
4. Easillyy, (n.d.) ‘a comprehensive guide to nftables’ in easillyy.com [online]. Available at: <http://easillyy.com/a-comprehensive-guide-to-nftables/> [accessed 26 February 2024]
5. Wiki, (2021) ‘What is nftable?’ in wiki.nftables.org [online]. Available at: [https://wiki.nftables.org/wiki-nftables/index.php/What\\_is\\_nftables%3F](https://wiki.nftables.org/wiki-nftables/index.php/What_is_nftables%3F) [accessed 27 February 2024]
6. LinuxCloudHacks (2023) Mastering Nftables Sets: A Comprehensive Guide Available at: <https://www.youtube.com/watch?v=YLVKuA4kiMA> [accessed 28 February 2024]
7. LinuxCloudHacks (2023) Protecting Incoming Traffic with Nftables Available at: [https://www.youtube.com/watch?v=K8JPwbcNy\\_0](https://www.youtube.com/watch?v=K8JPwbcNy_0) [accessed 28 February 2024]
8. LinuxCloudHacks (2023) Unleashing the Power of Nftables Chains and Verdict Maps Available at: <https://www.youtube.com/watch?v=Qy52m5hXiSg> [accessed 28 February 2024]
9. Wiki (2018) ‘Rate limiting matchings’ in wiki.nftables.org [online]. Available at: [https://wiki.nftables.org/wiki-nftables/index.php/Rate\\_limiting\\_matchings](https://wiki.nftables.org/wiki-nftables/index.php/Rate_limiting_matchings) [accessed 29 February 2024]
10. ZacsTech (2022) How to install Wireshark on Ubuntu 22.04 | 20.04 LTS Available at: <https://www.youtube.com/watch?v=IPxuiOLCtOY> [accessed 29 February 2024]